

ML4Cyber — Lab 4 API Reference

Pandas (pandas)

Docs: <https://pandas.pydata.org/docs/>

- Use `pandas.read_csv(...)` (https://pandas.pydata.org/docs/reference/api/pandas.read_csv.html) to load a CSV into a DataFrame.
- Use `df.copy()` (<https://pandas.pydata.org/docs/reference/api/pandas.DataFrame.copy.html>) to create a (deep) copy of a DataFrame before modifying it.
- Use `df.drop(...)` (<https://pandas.pydata.org/docs/reference/api/pandas.DataFrame.drop.html>) to remove columns or rows by label.
- Use `s.astype(...)` (<https://pandas.pydata.org/docs/reference/api/pandas.Series.astype.html>) to cast a Series to a new dtype (e.g., bytes → string).
- Use `s.value_counts()` (https://pandas.pydata.org/docs/reference/api/pandas.Series.value_counts.html) to count category frequencies.
- Use `df.values` (<https://pandas.pydata.org/docs/reference/api/pandas.DataFrame.values.html>) to get the underlying NumPy array for ML preprocessing.

NumPy (numpy)

Docs: <https://numpy.org/doc/stable/>

- Use `numpy.array(...)` (<https://numpy.org/doc/stable/reference/generated/numpy.array.html>) to create arrays (or to enforce an array type).
- Use `numpy.absolute(x) / numpy.abs(x)` (<https://numpy.org/doc/stable/reference/generated/numpy.absolute.html>) for elementwise absolute values.
- Use `numpy.nonzero(x)` (<https://numpy.org/devdocs/reference/generated/numpy.nonzero.html>) to get indices where elements are non-zero / True.
- Use `numpy.squeeze(x)` (<https://numpy.org/devdocs/reference/generated/numpy.squeeze.html>) to remove singleton dimensions.
- Use `x[:, None]` (a.k.a. `newaxis`) to add a singleton dimension for broadcasting; see NumPy indexing basics (https://numpy.org/devdocs/user/absolute_beginners.html).

scikit-learn (sklearn)

Docs: <https://scikit-learn.org/stable/>

- Call `sklearn.datasets.fetch_kddcup99(...)` (https://scikit-learn.org/stable/modules/generated/sklearn.datasets.fetch_kddcup99.html) to load the KDDCup'99 dataset.
- Call `train_test_split(X, y, ...)` (https://scikit-learn.org/stable/modules/generated/sklearn.model_selection.train_test_split.html) to create train/test partitions.

- Call `MinMaxScaler().fit_transform(X)` (<https://scikit-learn.org/stable/modules/generated/sklearn.preprocessing.MinMaxScaler.html>) to scale features into a fixed range (commonly $[0, 1]$).
- Call `StandardScaler().fit_transform(X)` (<https://scikit-learn.org/stable/modules/generated/sklearn.preprocessing.StandardScaler.html>) to standardize features (zero mean, unit variance).
- Call `PCA(n_components=...).fit_transform(X)` (<https://scikit-learn.org/stable/modules/generated/sklearn.decomposition.PCA.html>) for linear dimensionality reduction.
- Call `KernelPCA(n_components=..., kernel=...).fit_transform(X)` (<https://scikit-learn.org/stable/modules/generated/sklearn.decomposition.KernelPCA.html>) for nonlinear PCA via kernels.
- Call `IncrementalPCA(...).fit_transform(X)` (<https://scikit-learn.org/stable/modules/generated/sklearn.decomposition.IncrementalPCA.html>) for batch-wise PCA on larger datasets.
- Call `TSNE(n_components=2, ...).fit_transform(X)` (<https://scikit-learn.org/stable/modules/generated/sklearn.manifold.TSNE.html>) to produce a 2D embedding for visualization.

Matplotlib (`matplotlib`)

Docs: <https://matplotlib.org/stable/>

- Call `plt.figure(...)` (Pyplot tutorial: <https://matplotlib.org/stable/tutorials/pyplot.html>) to create a figure.
- Call `fig.add_subplot(...)` (https://matplotlib.org/stable/api/_as_gen/matplotlib.figure.Figure.add_subplot.html) to add axes (including 3D axes with `projection='3d'`).
- Call `ax.scatter(x, y, ...)` (https://matplotlib.org/stable/api/_as_gen/matplotlib.axes.Axes.scatter.html) for scatter plots.
- Call `plt.show()` (https://matplotlib.org/stable/api/_as_gen/matplotlib.pyplot.show.html) to display the figure.
- For 3D plotting, see `mplot3d` toolkit docs (<https://matplotlib.org/stable/api/toolkits/mplot3d.html>).

Python Standard Library

- Use `time.time()` and `time.sleep(...)` (time module docs: <https://docs.python.org/uk/3.8/library/time.html>) for simple timing and pacing.
- Use `@contextlib.contextmanager` (contextlib docs: <https://docs.python.org/3/library/contextlib.html>) to create lightweight context managers (e.g., for timing blocks).