

# ML4Cyber — Lab 5 API Reference

## Colab Setup (pip + dataset unpack)

Docs: <https://colab.research.google.com/>

- Use `%pip install numpy pandas scikit-learn matplotlib seaborn imbalanced-learn` to install required libraries.
- Use `!tar -xJf credit-card.tar.xz` to unpack the dataset archive in Colab.

## Pandas (pandas)

Docs: <https://pandas.pydata.org/docs/>

- Use `pandas.read_csv(...)` ([https://pandas.pydata.org/docs/reference/api/pandas.read\\_csv.html](https://pandas.pydata.org/docs/reference/api/pandas.read_csv.html)) to load `creditcard.csv` into a DataFrame.
- Use `df.head(n)` (<https://pandas.pydata.org/docs/reference/api/pandas.DataFrame.head.html>) to preview the first rows.
- Use `df.describe(...)` (<https://pandas.pydata.org/docs/reference/api/pandas.DataFrame.describe.html>) to get summary statistics (useful for percentiles).
- Use `df["col"].value_counts()` ([https://pandas.pydata.org/docs/reference/api/pandas.Series.value\\_counts.html](https://pandas.pydata.org/docs/reference/api/pandas.Series.value_counts.html)) to count fraud vs. non-fraud labels.
- Use `df.drop(..., axis=1)` (<https://pandas.pydata.org/docs/reference/api/pandas.DataFrame.drop.html>) to remove `Class` from features.
- Use boolean filtering like `df[df["Class"]==0]` to split fraud vs. non-fraud subsets.

## NumPy (numpy)

Docs: <https://numpy.org/doc/stable/>

- Use `np.arange(start, stop, step)` (<https://numpy.org/doc/stable/reference/generated/numpy.arange.html>) to create histogram bin edges.
- Use `np.arange(0.1, 1, 0.1)` (<https://numpy.org/doc/stable/reference/generated/numpy.arange.html>) to create percentile points for `describe(percentiles=...)`.
- Use `arr.reshape(-1, 1)` (<https://numpy.org/doc/stable/reference/generated/numpy.reshape.html>) to make a 2D column vector for scikit-learn scalers.

## Matplotlib (matplotlib.pyplot)

Docs: <https://matplotlib.org/stable/>

- Use `plt.hist(...)` ([https://matplotlib.org/stable/api/\\_as\\_gen/matplotlib.pyplot.hist.html](https://matplotlib.org/stable/api/_as_gen/matplotlib.pyplot.hist.html)) to plot class-conditional time distributions.

- Use `plt.hist(..., density=True)` to plot normalized distributions (fractions/density rather than raw counts).
- Use `alpha=...` in `plt.hist` to increase transparency when overlaying fraud vs. non-fraud.
- Use `plt.xlabel(...)` / `plt.ylabel(...)` ([https://matplotlib.org/stable/api/\\_as\\_gen/matplotlib.pyplot.xlabel.html](https://matplotlib.org/stable/api/_as_gen/matplotlib.pyplot.xlabel.html)) to label axes.
- Use `plt.title(...)` ([https://matplotlib.org/stable/api/\\_as\\_gen/matplotlib.pyplot.title.html](https://matplotlib.org/stable/api/_as_gen/matplotlib.pyplot.title.html)) to set plot titles.
- Use `plt.xlim((xmin, xmax))` ([https://matplotlib.org/stable/api/\\_as\\_gen/matplotlib.pyplot.xlim.html](https://matplotlib.org/stable/api/_as_gen/matplotlib.pyplot.xlim.html)) to constrain plot ranges.
- Use `plt.legend(...)` ([https://matplotlib.org/stable/api/\\_as\\_gen/matplotlib.pyplot.legend.html](https://matplotlib.org/stable/api/_as_gen/matplotlib.pyplot.legend.html)) to show series labels.
- Use `plt.show()` ([https://matplotlib.org/stable/api/\\_as\\_gen/matplotlib.pyplot.show.html](https://matplotlib.org/stable/api/_as_gen/matplotlib.pyplot.show.html)) to render the plot.

## Seaborn (`seaborn`)

Docs: <https://seaborn.pydata.org/api.html>

- Use `sns.histplot(...)` (<https://seaborn.pydata.org/generated/seaborn.histplot.html>) to plot smooth distributions.
- Use `kde=True` in `histplot` to overlay a KDE curve.
- Use `stat="density"` in `histplot` to plot a density (normalized) distribution.

## PCA Visualization (`sklearn.decomposition`)

Docs: <https://scikit-learn.org/stable/modules/decomposition.html>

- Use `IncrementalPCA(n_components=3)` (<https://scikit-learn.org/stable/modules/generated/sklearn.decomposition.IncrementalPCA.html>) to reduce V1-V28 to 3D.
- Use `ipca.fit_transform(X)` to compute the 3D embedding for plotting.
- If your lab provides helpers (e.g., `lab_6_util.visualize_samples`), use them to plot the 3D samples by class.

## Feature Scaling (`sklearn.preprocessing`)

Docs: <https://scikit-learn.org/stable/modules/preprocessing.html>

- Use `StandardScaler()` (<https://scikit-learn.org/stable/modules/generated/sklearn.preprocessing.StandardScaler.html>) to standardize Time after dividing by 24.
- Use `RobustScaler()` (<https://scikit-learn.org/stable/modules/generated/sklearn.preprocessing.RobustScaler.html>) to scale Amount robustly under outliers.
- Use `scaler.fit_transform(X)` to learn scaling parameters on training data and apply the transform.

- Caution: scalers expect 2D input; use `df[["Time"]].values` or `arr.reshape(-1,1)` and assign back via `df["Time"] = ...`

## Train/Test Split (`sklearn.model_selection`)

Docs: [https://scikit-learn.org/stable/modules/classes.html#module-sklearn.model\\_selection](https://scikit-learn.org/stable/modules/classes.html#module-sklearn.model_selection)

- Use `train_test_split(X, y, test_size=..., random_state=...)` ([https://scikit-learn.org/stable/modules/generated/sklearn.model\\_selection.train\\_test\\_split.html](https://scikit-learn.org/stable/modules/generated/sklearn.model_selection.train_test_split.html)) to create training and test sets.
- Use a larger `test_size` when the positive class is very rare, so the test set contains enough fraud cases.

## Imbalance Handling (`imblearn.over_sampling`)

Docs: [https://imbalanced-learn.readthedocs.io/en/stable/api.html#module-imblearn.over\\_sampling](https://imbalanced-learn.readthedocs.io/en/stable/api.html#module-imblearn.over_sampling)

- Use `SMOTE(random_state=...)` ([https://imbalanced-learn.readthedocs.io/en/stable/references/generated/imblearn.over\\_sampling.SMOTE.html](https://imbalanced-learn.readthedocs.io/en/stable/references/generated/imblearn.over_sampling.SMOTE.html)) to over-sample the minority class by synthetic interpolation.
- Use `BorderlineSMOTE(random_state=...)` ([https://imbalanced-learn.readthedocs.io/en/stable/references/generated/imblearn.over\\_sampling.BorderlineSMOTE.html](https://imbalanced-learn.readthedocs.io/en/stable/references/generated/imblearn.over_sampling.BorderlineSMOTE.html)) to over-sample mainly near the decision boundary.

## Pipelines (`imblearn.pipeline`)

Docs: <https://imbalanced-learn.readthedocs.io/en/stable/references/generated/imblearn.pipeline.Pipeline.html>

- Use `Pipeline(steps=[("smote", smote), ("dt", clf)])` to combine re-sampling + model training in one fit call.
- Use `pipeline.fit(X_train, y_train)` to apply SMOTE only on training data inside the pipeline.
- Use `pipeline.predict(X_test)` to run the fitted pipeline on the untouched test set.

## Decision Trees (`sklearn.tree`)

Docs: <https://scikit-learn.org/stable/modules/tree.html>

- Use `DecisionTreeClassifier(...)` (<https://scikit-learn.org/stable/modules/generated/sklearn.tree.DecisionTreeClassifier.html>) to create a decision tree classifier.
- Use `dt.fit(X_train, y_train)` to train on original (non-resampled) data.
- Use `dt.predict(X_test)` to produce predictions for evaluation.

## Evaluation (`sklearn.metrics`)

Docs: <https://scikit-learn.org/stable/modules/classes.html#module-sklearn.metrics>

- Use `classification_report(y_true, y_pred)` ([https://scikit-learn.org/stable/modules/generated/sklearn.metrics.classification\\_report.html](https://scikit-learn.org/stable/modules/generated/sklearn.metrics.classification_report.html)) to print precision/recall/F1 per class.
- Use `confusion_matrix(y_true, y_pred)` ([https://scikit-learn.org/stable/modules/generated/sklearn.metrics.confusion\\_matrix.html](https://scikit-learn.org/stable/modules/generated/sklearn.metrics.confusion_matrix.html)) to compute TP/FP/FN/TN counts.
- Use precision/recall/F1 (not just accuracy) when the dataset is highly imbalanced.

## Common TODO Patterns in This Lab

Docs: [https://matplotlib.org/stable/api/\\_as\\_gen/matplotlib.pyplot.hist.html](https://matplotlib.org/stable/api/_as_gen/matplotlib.pyplot.hist.html)

- Time-by-class histogram: use `plt.hist(time_nonfraud, bins, density=True, alpha=..., label="Non-Fraud")` and a second `plt.hist` for fraud with a different color/label.
- Time conversion: use `df["Time"]/(60*60)` to convert seconds to hours; use `time/24` to convert hour-of-day ratio before `StandardScaler`.
- Scaling columns: assign back to DataFrame columns after `fit_transform` (remember scalers return 2D arrays).
- SMOTE + DT training: prefer a `Pipeline` so over-sampling happens only on the training split.