

EE 467 — Lab 6 API Reference

Ensemble Learning and Random Forest

Colab Setup (pip + dataset unpack)

Docs: <https://colab.research.google.com/>

- Use `%pip install numpy pandas scikit-learn matplotlib seaborn` to install required libraries (if not already installed).
- Use `!tar -xf credit-card.tar.xz` to unpack the dataset archive in Colab.
- Keep `creditcard.csv` in the current working directory so `pd.read_csv("./creditcard.csv")` works.

Pandas (pandas)

Docs: <https://pandas.pydata.org/docs/>

- Use `pandas.read_csv(...)` (https://pandas.pydata.org/docs/reference/api/pandas.read_csv.html) to load `creditcard.csv`.
- Use `df.drop(["Class"], axis=1).values` (<https://pandas.pydata.org/docs/reference/api/pandas.DataFrame.drop.html>) to create the feature matrix.
- Use `df["Class"].values` to create the label vector (fraud vs. non-fraud).

Train/Test Split (sklearn.model_selection)

Docs: https://scikit-learn.org/stable/modules/classes.html#module-sklearn.model_selection

- Use `train_test_split(X, y, test_size=..., random_state=...)` (https://scikit-learn.org/stable/modules/generated/sklearn.model_selection.train_test_split.html) to create `feat_train`, `feat_test`, `y_train`, `y_test`.
- Caution: keep `random_state` fixed so results are reproducible across runs.

Feature Scaling (sklearn.preprocessing)

Docs: <https://scikit-learn.org/stable/modules/preprocessing.html>

- Use `StandardScaler()` (<https://scikit-learn.org/stable/modules/generated/sklearn.preprocessing.StandardScaler.html>) to scale Time after converting seconds to hour-of-day.
- Use `RobustScaler()` (<https://scikit-learn.org/stable/modules/generated/sklearn.preprocessing.RobustScaler.html>) to scale Amount (robust to outliers).
- Use `scaler.fit_transform(arr[:, None])` to scale a single column (scalers expect 2D input).
- Caution: in real ML workflows, fit scalers on **training** only; this lab follows the provided starter code for consistency.

Utility Functions (`lab_6_util`)

Docs: (provided in the lab starter / local module)

- Use `timeit(...)` as a context manager to time code blocks.
- Use `evaluate_model(model, name, X_test, y_test)` to print accuracy, precision, recall, and F1-score.
- Caution: only call `evaluate_model` on a **trained** model (after `fit`).

Baseline Model: Logistic Regression (`sklearn.linear_model`)

Docs: https://scikit-learn.org/stable/modules/linear_model.html

- Use `LogisticRegression(max_iter=200)` (https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.LogisticRegression.html) to avoid non-convergence warnings on large datasets.
- Use `model.fit(X_train, y_train)` then `evaluate_model(...)`.

Voting Ensemble (`sklearn.ensemble.VotingClassifier`)

Docs: <https://scikit-learn.org/stable/modules/ensemble.html>

- Use `VotingClassifier(estimators=[...], voting="hard")` (<https://scikit-learn.org/stable/modules/generated/sklearn.ensemble.VotingClassifier.html>) to build a majority-vote ensemble.
- For the sub-classifiers, use:
 - `LogisticRegression(max_iter=200)` (docs above),
 - `GaussianNB()` (https://scikit-learn.org/stable/modules/generated/sklearn.naive_bayes.GaussianNB.html),
 - `DecisionTreeClassifier()` (<https://scikit-learn.org/stable/modules/generated/sklearn.tree.DecisionTreeClassifier.html>).
- Use `voting_ensemble.fit(X_train, y_train)` then `evaluate_model(...)`.
- Use `voting_ensemble.named_estimators_` to access trained sub-models by name and evaluate each one.
- Soft voting: set `voting="soft"` (requires sub-models that support `predict_proba`).

Bagging (`sklearn.ensemble.BaggingClassifier`)

Docs: <https://scikit-learn.org/stable/modules/ensemble.html#bagging>

- Use `BaggingClassifier(base_estimator=..., n_estimators=..., max_samples=...)` (<https://scikit-learn.org/stable/modules/generated/sklearn.ensemble.BaggingClassifier.html>) to train many copies of the same base estimator on random subsets.
- Typical knobs used in this lab:

- `n_estimators`: number of sub-classifiers,
- `max_samples`: fraction (or count) of training samples per classifier,
- `bootstrap`: whether to sample with replacement,
- `bootstrap_features + max_features`: sub-sample features per classifier.
- Common pattern: create a dictionary `bagging_models = { "setting": model, ... }` and loop over it to `fit` + `evaluate_model`.

Random Forest (`sklearn.ensemble.RandomForestClassifier`)

Docs: <https://scikit-learn.org/stable/modules/ensemble.html#random-forests>

- Use `RandomForestClassifier(n_estimators=...)` (<https://scikit-learn.org/stable/modules/generated/sklearn.ensemble.RandomForestClassifier.html>) to train an ensemble of decision trees (bagging + feature randomness).
- Use `n_jobs=...` to train trees in parallel (speed-up on multi-core machines).
- If the starter defines `N_ENSEMBLE_CPUS = max(os.cpu_count()//2, 1)`, pass `n_jobs=N_ENSEMBLE_CPUS`.

Boosting (`sklearn.ensemble`)

Docs: <https://scikit-learn.org/stable/modules/ensemble.html#boosting>

- Use `AdaBoostClassifier()` (<https://scikit-learn.org/stable/modules/generated/sklearn.ensemble.AdaBoostClassifier.html>) for adaptive boosting (re-weights samples over rounds).
- Use `GradientBoostingClassifier(n_estimators=...)` (<https://scikit-learn.org/stable/modules/generated/sklearn.ensemble.GradientBoostingClassifier.html>) for gradient boosting (sequentially fits learners to correct errors).
- Use `fit` then `evaluate_model` to compare boosting models to other ensembles.

Stacking (`sklearn.ensemble.StackingClassifier`)

Docs: <https://scikit-learn.org/stable/modules/ensemble.html#stacking>

- Use `StackingClassifier(estimators=[...], final_estimator=..., n_jobs=...)` (<https://scikit-learn.org/stable/modules/generated/sklearn.ensemble.StackingClassifier.html>) to build a meta-model that learns from base-model predictions.
- Base models used in this lab often include:
 - `RandomForestClassifier(n_estimators=40)` (docs above),
 - `LogisticRegression(max_iter=200)` (docs above),
 - `LinearSVC(max_iter=1500)` (<https://scikit-learn.org/stable/modules/generated/sklearn.svm.LinearSVC.html>).
- Meta-model (final estimator): commonly `LogisticRegression()`.
- Caution: some models (e.g., `LinearSVC`) do not provide `predict_proba`; stacking can use alternative scores internally (per scikit-learn behavior).

Evaluation (`sklearn.metrics`)

Docs: <https://scikit-learn.org/stable/modules/classes.html#module-sklearn.metrics>

- Use `classification_report(y_true, y_pred)` (https://scikit-learn.org/stable/modules/generated/sklearn.metrics.classification_report.html) to print precision/recall/F1 per class.
- Use `confusion_matrix(y_true, y_pred)` (https://scikit-learn.org/stable/modules/generated/sklearn.metrics.confusion_matrix.html) to inspect TP/FP/FN/TN counts.
- For highly imbalanced fraud detection, focus on **precision/recall/F1** (accuracy alone can be misleading).

Common TODO Patterns in This Lab

Docs: <https://scikit-learn.org/stable/modules/ensemble.html>

- Voting ensemble: create `VotingClassifier` with three named estimators; `fit` once; evaluate ensemble, then evaluate each sub-model via `named_estimators_`.
- Soft voting: switch `voting="hard"` to `voting="soft"` and re-evaluate (requires probability outputs).
- Bagging grid: store multiple `BaggingClassifier` configurations in a dictionary and loop over `fit + evaluate_model`.
- Random forest: train two models with different `n_estimators` (e.g., 40 vs 100); use `n_jobs=N_ENSEMBLE_CPUS` to reduce training time.
- Stacking: define base estimators list + meta-estimator; call `fit` once; evaluate the final stacking ensemble on `feat_test`.